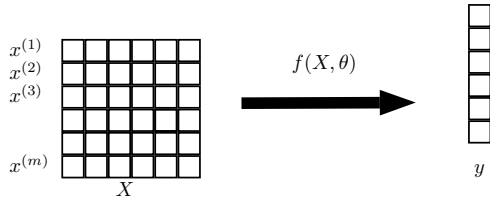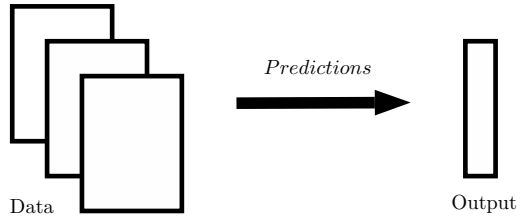POLIMI GRADUATE SCHOOL OF MANAGEMENT

MACHINE LEARNING IN PYTHON - LAB SESSION

Andrea Mor - andrea.mor@polimi.it - https://github.com/mrondr/BABD2023
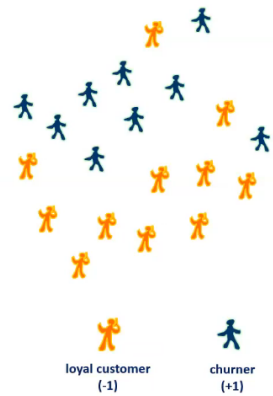
SUPERVISED LEARNING

Data $\xrightarrow{Predictions}$ Output

$$X \xrightarrow{f(X, \theta)} y$$

where X has rows $x^{(1)}, x^{(2)}, x^{(3)}, \ldots, x^{(m)}$

THE PROBLEM: BANK TELEMARKETING[1]

Attribute		Type	Description/Values
Personal	age	num	Age of the potential client
	job	cat	admin., blue- collar, entrepreneur, housemaid, … , unknown
	marital_status	cat	divorced, married, single, unknown
	education	cat	basic.4y, basic.6y, basic.9y, high.school, … , unknown
Bank	default	cat	The client has credit in default: no,yes,unknown
	housing	cat	The client has a housing loan contract: no,yes,unknown
	loan	cat	The client has a personal loan: no,yes,unknown
Campain	contact	cat	Communication type: cellular,telephone
	month	cat	Last month contacted: jan, feb ,…, dec
	day_of_week	cat	Last contact day : mon, tue,…, fri
	duration	num	Last contact duration (in seconds)
	campain	num	Number of contacts performed during this campaign
	pdays	num	Number of days that passed by after last contact
	previous	num	Number of contacts performed before this campaign
	poutcome	cat	Outcome of the prev. marketing campaign: failure,nonexistent,success
Economical	emp.var.rate	num	Employment variation rate in the last quarter
	cons.price.idx	num	Consumer price index in the last month
	cons.conf.idx	num	Monthly consumer confidence index
	euribor3m	num	Dayly Euro Interbank Offered Rate
	nr.employed	num	Number of employed citizens in the last quarter (thousands)
Target	**success**	**target**	**0: no, 1: yes**

[1] A data-driven approach to predict the success of bank telemarketing. S. Moroa, P. Cortez, P. Rita. Decision Support Systems, 62:22-31, 2014.
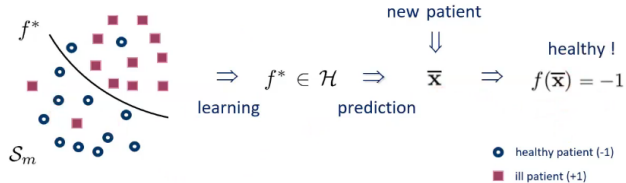
CLASSIFICATION PROBLEM

CLASSIFICATION FORMULATION

$S_m = \{(\mathbf{x}_i, y_i), \ i \in \mathcal{M}\}$: **training set, where** $\mathbf{x}_i \in \Re^n$ **and** $y_i \in \mathcal{D}$

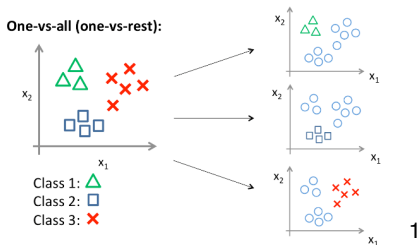$\mathcal{H}$ **denotes a set of functions** $f(\mathbf{x}) : \Re^n \longmapsto \mathcal{D}$

Classification problem: define a hypotheses space $\mathcal{H}$ **and a function** $f^* \in \mathcal{H}$
which optimally describes the relationship between $\mathbf{x}_i$ **and** y_i

$$\Rightarrow \quad f^* \in \mathcal{H} \quad \Rightarrow \quad \overline{\mathbf{x}} \quad \Rightarrow \quad f(\overline{\mathbf{x}}) = -1$$

learning   prediction

new patient ⇓   healthy !

○ healthy patient (-1)
■ ill patient (+1)

MULTI-CLASS CLASSIFICATION

1. **One-vs-Rest** We perform $|H|$ different binary classifications: one for every class.

We decide based on a majority vote.

2. **One-vs-One** We perform $|H|(|H|-1|)/2$ binary classifications: one for every pair of classes. We decide based on a majority vote.

Image via www.cc.gatech.edu

CLASSIFICATION MODELS

► Heuristics Methods
  - Nearest Neighbours
  - Classification Trees

► Probabilistic Methods
  - Bayesian Methods

► Regression Methods
  - Logistic regression

► Separation Methods
  - Support vector machine
  - Perceptron
  - Neural Networks

EVALUATION DIMENSIONS

- ► Prediction accuracy
- ► Speed
- ► Robustness
- ► Scalability
- ► Intepretability
- ► Rules effectiveness

QUALITY MEASURES - CONFUSION MATRIX

Prediction outcome

	0	1
Actual value 0	True Negative	False Positive
1	False Negative	True Positive

- **Accuracy** = $\frac{TP+TN}{TP+FP+TN+FN}$
  When: balanced problems, equal costs of FN and FP.
  When NOT: imbalanced problems.

- **Recall** (True Positive rate)= $\frac{TP}{FN+TP}$
  "proportion of true positives among actual positive"
  When: you want to be sure to catch all the 1s, at the cost of having false alerts.

- **Precision** (Positive Predictive value) = $\frac{TP}{TP+FP}$
  "proportion of true positives among positive predictions"
  When: you want to be sure that all that you predict as 1, are indeed 1, e.g., false alerts are costly.

- Precision and Recall are typically evaluated together.

QUALITY MEASURES - CONFUSION MATRIX

Prediction outcome

	0	**1**
0	True Negative	False Positive
1	False Negative	True Positive

Actual value

▶ **Geom. mean**=$\sqrt{\text{Precision} \times \text{Recall}}$

▶ **F-score**=$\frac{(1+\beta^2)}{\beta^2}\frac{1}{\frac{1}{\text{Precision}} + \frac{1}{\text{Recall}}} = (1+\beta^2)\frac{\text{Precision·Recall}}{(\beta^2 \cdot \text{Precision}) \cdot \text{Recall}}$

▶ **F1**: harmonic mean of Prec. and Recall
$= \frac{2TP}{2TP+FP+FN}$
Frequently the to-go metric.

▶ **F2**: 2x emphasis on Recall
When: False Negatives are more costly.

QUALITY MEASURES - CONFUSION MATRIX (AUXILLIARY)

Prediction outcome

	0	**1**
0	True Negative	False Positive
1	False Negative	True Positive

Actual value

- ► **False Positive rate** = $\frac{FP}{FP+TN}$
  "proportion of false positives among actual negatives"
  When: cost of dealing with positive prediction is high (e.g., false alarms).

- ► **True Negative rate** (Specificity) = $\frac{TN}{TN+FP}$
  When: e.g., you want to be sure when you say something is safe for health.

QUALITY MEASURES - ROC CURVE & AUC
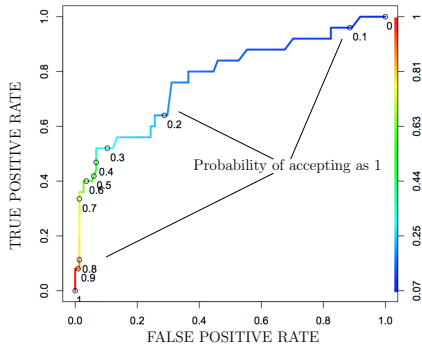
- ▶ If we accepting even with small probability then $TPR = FPR = 1$

- ▶ If we accepting just with high probability then $TPR = FPR = 0$

- ▶ The perfect classificator is the the point $(0, 1)$

- ▶ $AUC \in [0.5, 1]$ area under the curve is a quality measure of our algorithm.
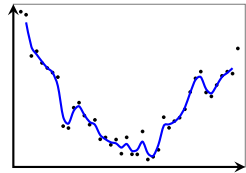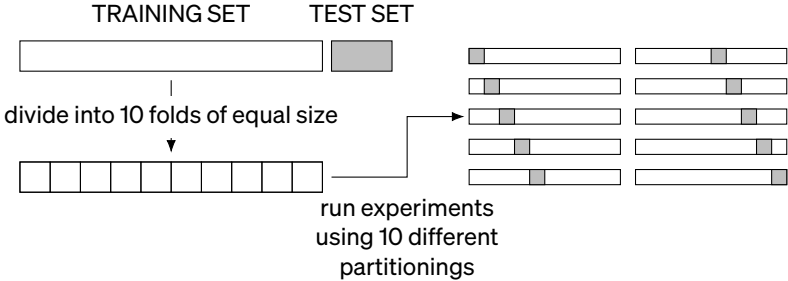
WORKFLOW

1. Data Exploration/Analysis

Data Exploration
Data Preprocessing
- Missing/Inconsistent data
- Noisy data
- Conversion (categorical)

Previous Data

2. Model Creation

Training Set

Algorithm Selection

Training
Hyper-parameter selection

Evaluation

Test Set

3. Predictions

New Data

Data Preprocessing

Model

Predictions
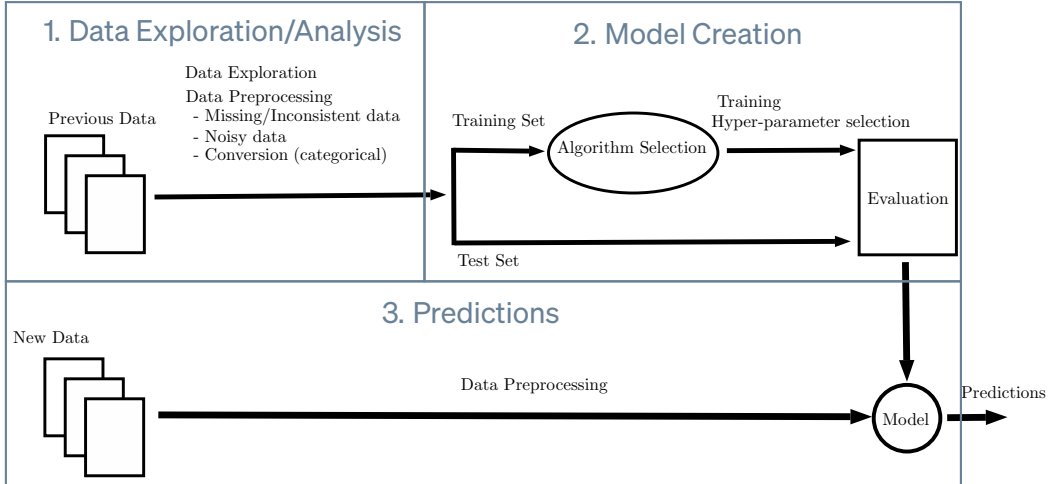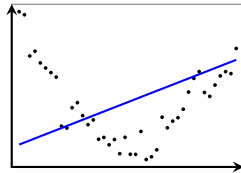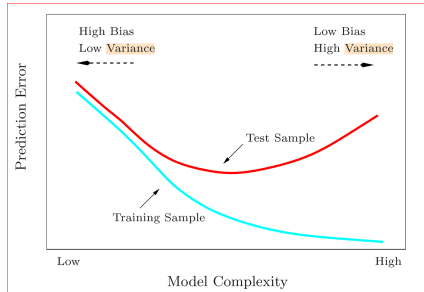
UNDER/OVER-FITTING
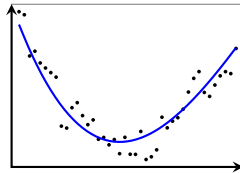
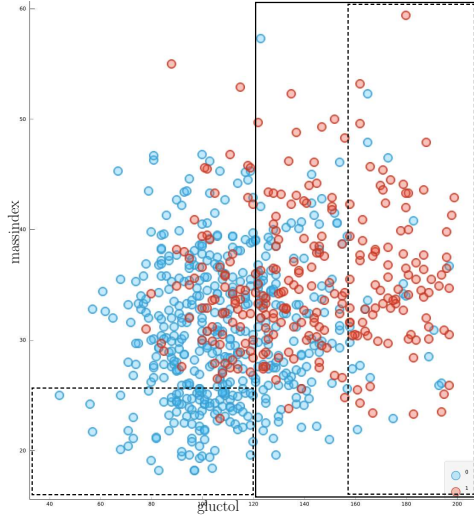Underfitting   Balance   Overfitting

CROSS VALIDATION

TRAINING SET   TEST SET

divide into 10 folds of equal size

run experiments
using 10 different
partitionings

POLIMI GRADUATE SCHOOL OF MANAGEMENT

THE ALGORITHMS

K-NEAREST NEIGHBORS
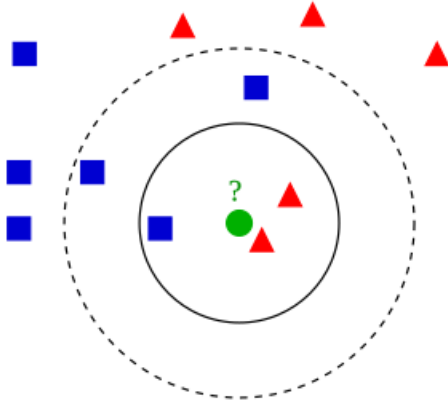
Main Parameters

- ► k : number of neighbors
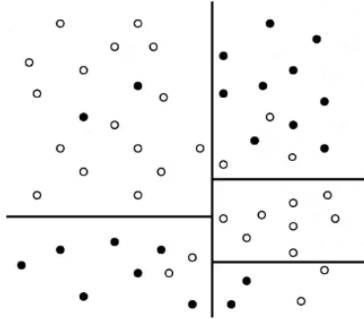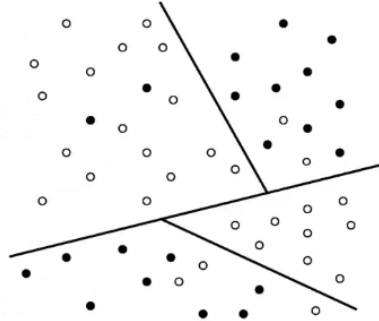- ► neighbor weights
- ► distances

CLASSIFICATION TREE
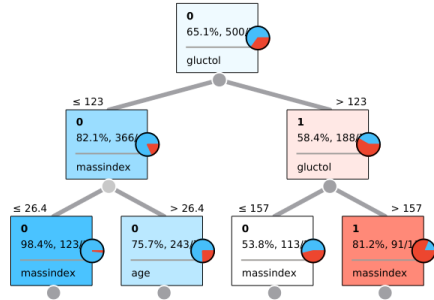
Tree

types

- ► Binary tree (zero/two descendants)
- ► General trees
- ► Uni-variate tree ($X_j < b$)
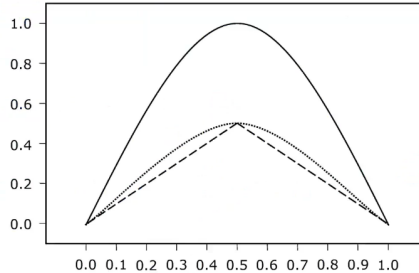- ► Multi-variate tree ($\sum_{j=1}^{n} w_j x_j = b$)
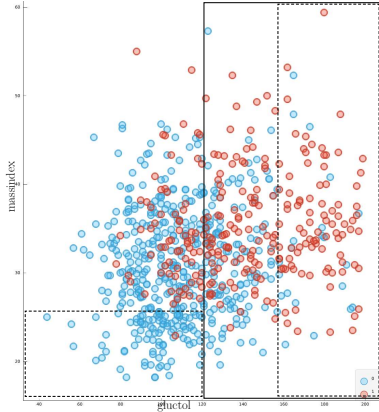
CLASSIFICATION TREE

classification by an
axis parallel tree

classification by an
oblique tree

CLASSIFICATION TREE

- ► **Entropy index:**
- ► **Gini index:**
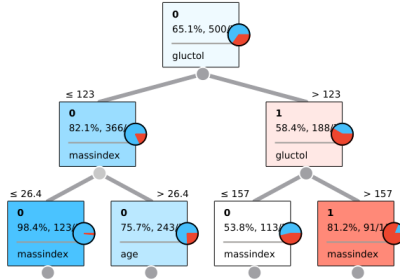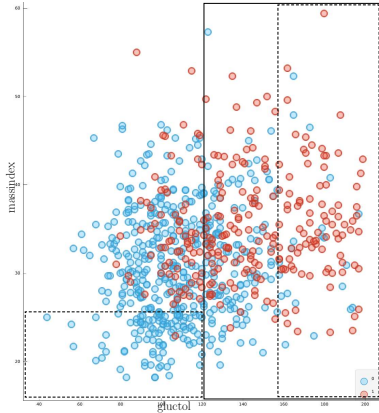- ► **Miss-classification index:**

Split criteria

$$-\sum_{h=1}^{H} f_h \log_2 f_h$$

$$1 - \sum_{h=1}^{H} f_h^2$$

$$1 - \max_h f_h$$

CLASSIFICATION TREE

Main Parameters

▶ impurity measure: "gini", "entropy"

▶ max_depth

▶ min_samples_split: minimum number of samples to split an internal node

▶ min_sample_leaf: minimum number of samples required to be at a leaf node

NAIVE BAYESIAN CLASSIFIER

▶ Bayes Theorem

$$P(y|\mathbf{x}) = \frac{P(\mathbf{x}|y)P(y)}{P(\mathbf{x})} = \frac{P(\mathbf{x}|y)\,P(y)}{\sum_{l=1}^{H} P(\mathbf{x}|y)\,P(y)}$$

▶ Maximum a posteriori hypothesis

$$y_{MAP} = \arg\max_{y \in \mathcal{H}} P(y|\mathbf{x}) = \arg\max \frac{P(\mathbf{x}|y)P(y)}{P(\mathbf{x})}$$

▶ Independence (Naive)

$$P(\mathbf{x}|y) = P(x_1|y) \cdot P(x_2|y) \cdot \ldots \cdot P(x_n|y) = \prod_{j=1}^{n} P(x_j|y)$$

NAIVE BAYESIAN CLASSIFIER

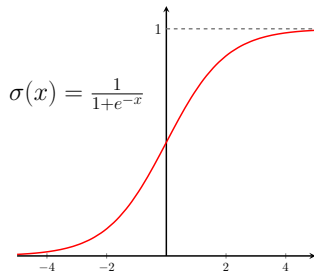▶ Categorical/discrete attributes

$$P(x_j|y) = P(x_j = r_{jk}|y = v_h)$$

$\rightarrow$ empirical frequency of the observed value on the class v_h

▶ Numerical attribute

$$P(x_j|y) \sim N(\mu_{jh}, \sigma_{jh})$$

$\rightarrow$ assuming Gaussian density with empirical parameters

LOGISTIC REGRESSION

The model postulates that

$$\log\left[\frac{P(y=1|x)}{P(y=0|x)}\right] = w^T x$$

$\sigma(x) = \frac{1}{1+e^{-x}}$

then if $p = P(y=1|x)$:

$$\frac{p}{1-p} = e^{w^T x}$$

$$\Rightarrow p = P(y=1|x) = \frac{e^{w^T x}}{1+e^{w^T x}} = \frac{1}{1+e^{-w^T x}},$$

$$P(y=0|x) = 1 - p = \frac{1}{1+e^{w^T x}} = \frac{e^{-w^T x}}{1+e^{-w^T x}}$$

LOGISTIC REGRESSION

Therefore if we maximizes the likelihood

$$L(w) := P(Y_1 = y_1, \ldots, Y_m = y_m | w, x_1, \ldots, x_m) = \Pi_{i=1}^n P(Y_i = y_i | w, x_1, \ldots, x_m).$$

Assuming independence:

$$L(w) := \prod_{i | y_i = 1} p(x_i) \cdot \prod_{i | y_i = 0} (1 - p(x_i))$$

$$L(w) := \prod_{i=1}^n p(x_i)^{y_i} (1 - p(x_i))^{1 - y_i}$$

is equivalent to maximize the log likelihood

$$l(w) = \sum_{i=1}^n (y_i log(p(x_i)) + (1 - y_i) log(1 - p(x_i)))$$

$$= \sum_{i=1}^n (y_i log(\frac{p(x_i)}{1 - p(x_i)}) + log(1 - p(x_i)))$$

Given the quantities defined above

$$l(w) = \sum_{i=1}^n (y_i w^T x_i) - log(1 + e^{-w^T x_i}))$$

Therefore we aim to resolve the following optimization problem

$$\max_w l(w)$$

LOGISTIC REGRESSION

Regulation term(s) can be considered (recall that max $l(w)$ is equivalent to min $-l(w)$):

$$\min_w \frac{1}{2}||w||^2 - C \cdot l(w)$$

Main Parameters

- ▶ C: Inverse of regularization strength
- ▶ Resolution algorithm parameters:
  - solver: lbfgs, newton-cg, liblinear, sag, saga.
  - tol: Tolerance for stopping criteria.
  - max_iter: max. number of iterations
  - n_jobs: Number of CPU cores

MULTI-LAYER PERCEPTRON - SMALL EXAMPLE

Temp. [C]	20	31	15	18	21
Humidity [%]	40	36	23	45	30
Prob. Rain	0.70	0.52	0.55	0.73	0.60

MULTI-LAYER PERCEPTRON - SMALL EXAMPLE

Temp. [C]	20	31	15	18	21
Humidity [%]	40	36	23	45	30
Prob. Rain	0.70	0.52	0.55	0.73	0.60

Temperature $\longrightarrow$ T $\quad w_1$

Humidity $\longrightarrow$ H $\quad w_2$ $\longrightarrow$ f $\longrightarrow$ Output
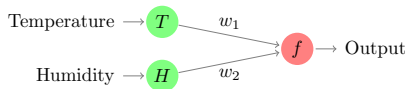
$$f(T, H, w_1, w_2) = \underbrace{\sigma}_{\text{activation function}} (w_1 \cdot T + w_2 \cdot H)$$

$$\sigma(x) = \frac{1}{1+e^{-x}}$$

$$\min_{w_1, w_2} \sum_{i=1}^{5} [P_i(\text{rain}) - f(T_i, H_i, w_1, w_2)]^2$$
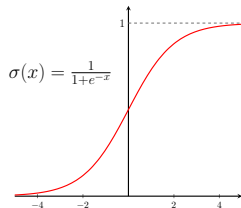
For a classification problem we can use the Likelihood as cost function.

MULTI-LAYER PERCEPTRON - SMALL EXAMPLE

$$\min_{w_1,w_2} \sum_{i=1}^{5} [P_i(\mathsf{rain}) - f(T_i, H_i, w_1, w_2)]^2$$

$$= \mathsf{min} \left[0.7 - 1/(1 + e^{-(w_1 \cdot 20 + w_2 \cdot 0.4)})\right]^2 + \left[0.52 - 1/(1 + e^{-(31 \cdot w_1 + w_2 \cdot 0.36)})\right]^2 + \cdots$$

MULTI-LAYER PERCEPTRON - SMALL EXAMPLE

$$\min_{w_1, w_2} \sum_{i=1}^{5} [P_i(\text{rain}) - f(T_i, H_i, w_1, w_2)]^2$$

$$= \min \left[0.7 - 1/(1 + e^{-(w_1 \cdot 20 + w_2 \cdot 0.4)}) \right]^2 + \left[0.52 - 1/(1 + e^{-(31 \cdot w_1 + w_2 \cdot 0.36)}) \right]^2 + \cdots$$
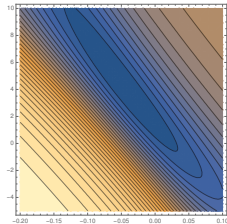
$$(w_1^*, w_2^*) = (-0.044, 4.147)$$

MULTI-LAYER PERCEPTRON – SMALL EXAMPLE

$$\min_{w_1, w_2} \sum_{i=1}^{5} [P_i(\text{rain}) - f(T_i, H_i, w_1, w_2)]^2$$

$$= \min \left[0.7 - 1/(1 + e^{-(w_1 \cdot 20 + w_2 \cdot 0.4)})\right]^2 + \left[0.52 - 1/(1 + e^{-(31 \cdot w_1 + w_2 \cdot 0.36)})\right]^2 + \cdots$$
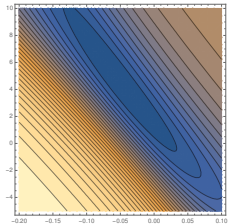
$$(w_1^*, w_2^*) = (-0.044, 4.147)$$

Temp. [C]	20	31	15	18	21
Humidity [%]	40	36	23	45	30
Prob. Rain	0.70	0.52	0.55	0.73	0.60
Predicted	0.70	0.56	0.58	0.75	0.60
Error	0.0	-0.04	-0.03	-0.02	0.0

MULTI-LAYER PERCEPTRON
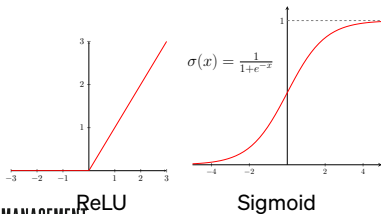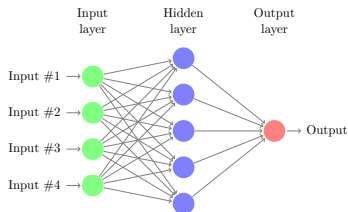
Main Parameters

- ► hidden_layer_sizes: $(n_1, n_2, \ldots, n_L)$
- ► activation: identity, logistic, tanh, relu
- ► alpha regularization term parameter
- ► Resolution algorithm parameters: solver, tol, batch_size, learning_rate, max_iter.
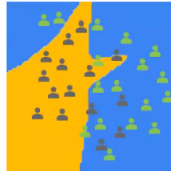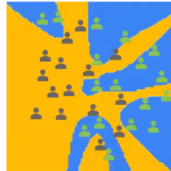
NEURAL NETWORK
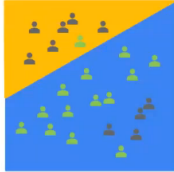
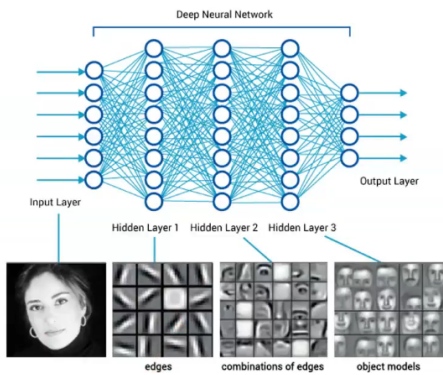1 LAYER   2 LAYER   3+ LAYER
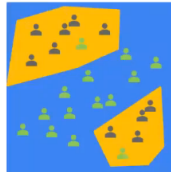
1 LAYER, 3 NODES   1 LAYER, 6 NODES   1 LAYER, 20 NODES

DEEP LEARNING

Artificial Neural Network

Input Layer

Hidden Layer 1

Output Layer

Deep Neural Network

Input Layer

Hidden Layer 1   Hidden Layer 2   Hidden Layer 3

Output Layer

edges   combinations of edges   object models
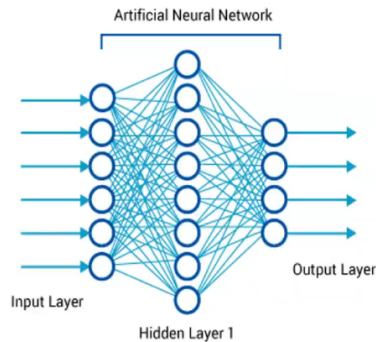
CLASSIFICATION AND STATISTICAL LEARNING THEORY

Given an hypothesis space $\mathcal{F}$, a predictive function $f \in \mathcal{F}$, and a loss function $V(y, f(\mathbf{x}))$, we define

▶ the **empirical risk** on a train set $\mathcal{T}$ as

$$R_{emp}(f) = \frac{1}{m} \sum_{i=1}^{m} V(y_i, f(\mathbf{x_i}))$$

▶ and the **expected risk** as

$$R(f) = \frac{1}{2} \int V(y, f(\mathbf{x})) \, dP(\mathbf{x}, y)$$

Overfitting arises when the difference between these errors is large.

VC-DIMENSION

The Vapnik–Chervonenkis (VC) dimension is the maximum number of points that can be correctly classified by classifiers in $\mathcal{F}$. We can prove that

$$R(f) - R_{emp}(f) \leq \sqrt{\tfrac{1}{t}(\gamma \log(2t/\gamma) - \log(\eta/4) + 1)} = \Psi(t, \gamma, \eta)$$
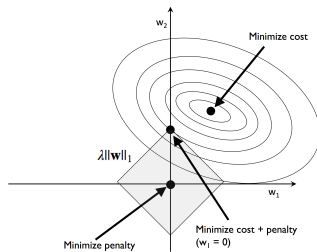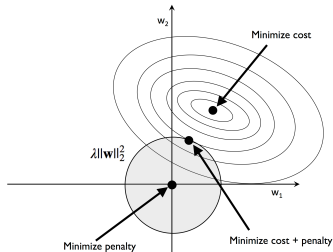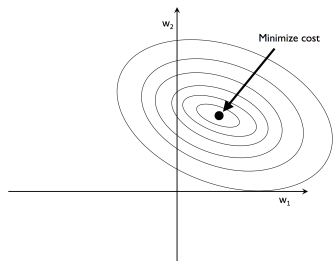
with probability $1 - \eta$, where

- ▶ t is the number of training points
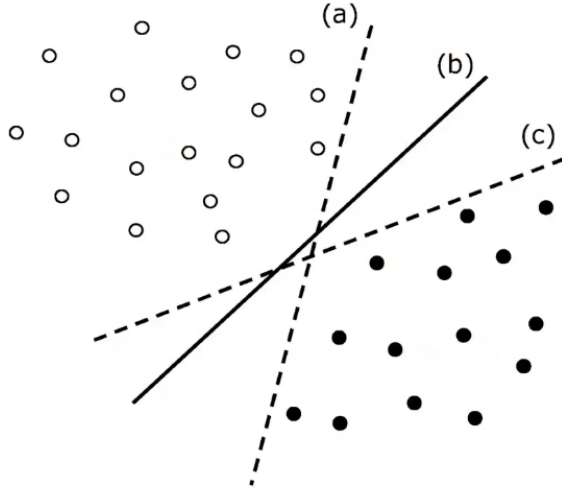- ▶ γ is the VC-dimension
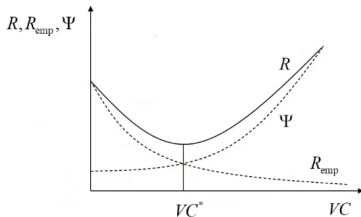
REGULARIZATION
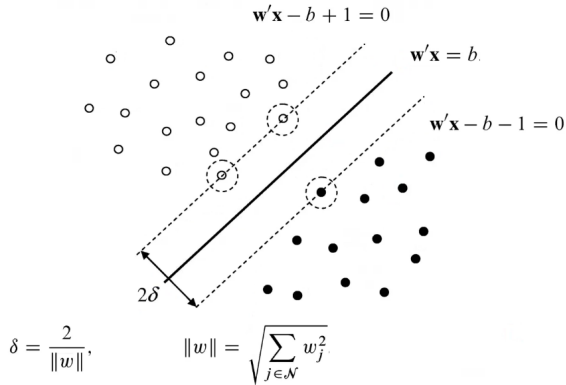
Structural risk minimization (SRM)

$$\hat{R}(f) = \frac{1}{t} \sum_{i=1}^{t} V(y_i, f(\mathbf{x_i})) + \lambda \|f\|_K^2$$
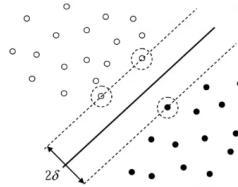
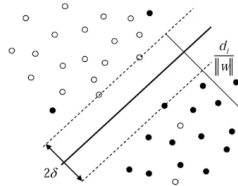# SUPPORT VECTOR MACHINE - LINEARLY SEPARABLE

SVM – LINEARLY SEPARABLE

$$\mathbf{w}'\mathbf{x} - b + 1 = 0$$

$$\mathbf{w}'\mathbf{x} = b$$

$$\mathbf{w}'\mathbf{x} - b - 1 = 0$$

$$2\delta$$

$$\delta = \frac{2}{\|w\|}, \qquad \|w\| = \sqrt{\sum_{j \in \mathcal{N}} w_j^2}$$

$$\min_{\mathbf{w}, b} \quad \frac{1}{2} \|\mathbf{w}\|_2^2$$

$$\text{s. t.} \quad y_i(\mathbf{w}'\mathbf{x}_i - b) \geq 1 \quad i \in \mathcal{M}$$
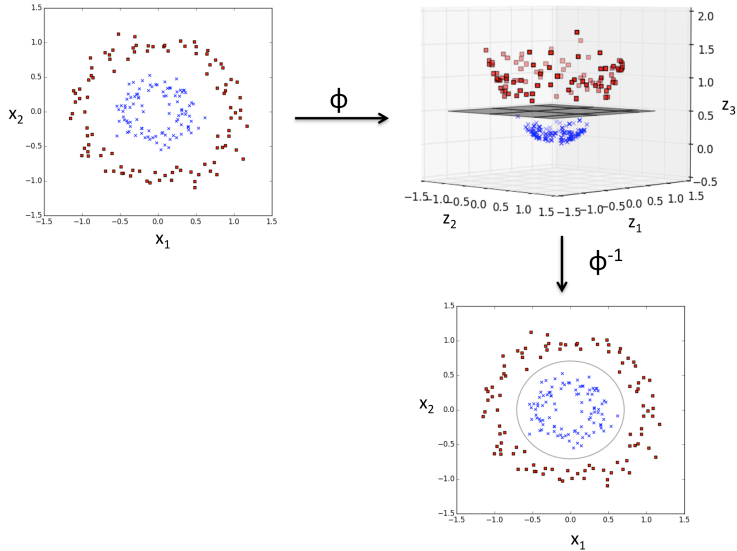
SVM - GENERAL CASE

$$\min_{\mathbf{w},b} \quad \frac{1}{2}\|\mathbf{w}\|_2^2$$
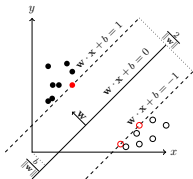$$\text{s.t.} \quad y_i(\mathbf{w}'\mathbf{x}_i - b) \geq 1 \quad i \in \mathcal{M}$$

$$\min_{\mathbf{w},b,d} \quad \frac{1}{2}\|\mathbf{w}\|_2^2 + \lambda \sum_{i=1}^{m} d_i$$
$$\text{s.t.} \quad y_i(\mathbf{w}'\mathbf{x}_i - b) \geq 1 - d_i \quad i \in \mathcal{M}$$
$$\quad d_i \geq 0 \quad i \in \mathcal{M}$$

SVM - KERNELS

SVM - GENERAL CASE

SVC with linear kernel

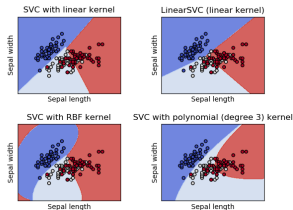LinearSVC (linear kernel)

SVC with RBF kernel

SVC with polynomial (degree 3) kernel

$$\min_{w,b,d} \quad \frac{1}{2}||w||^2 + C \sum_{i=1}^{m} d_i$$

$$\text{subject to } y_i(w^T \underbrace{\phi(x_i)}_{\text{kernel}} - b) \geq 1 - d_i,$$
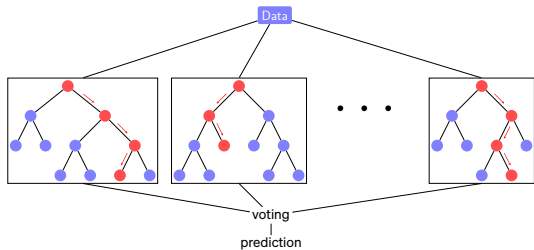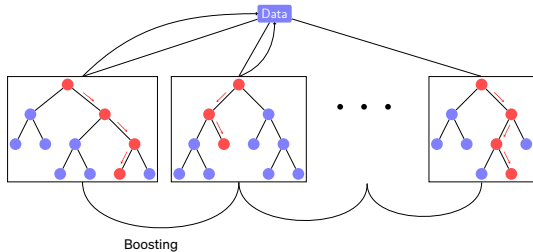
$$d_i \geq 0$$

Main Parameters

▶ C: Inverse of regularization strength

▶ kernel:
- linear: $x'x$
- poly: $(\gamma x'x + r)^d$
- rbf: $exp(-\gamma||x - x'||^2)$
- sigmoid: $\tanh(\gamma x'x + r)$

▶ degree(d), gamma(γ), coef0(r)

▶ Resolution algorithm parameters

ENSEMBLE METHODS

Bagging

Boosting

voting
prediction

Boosting

RANDOM FOREST

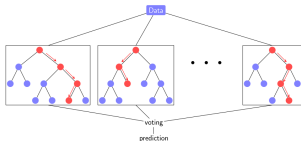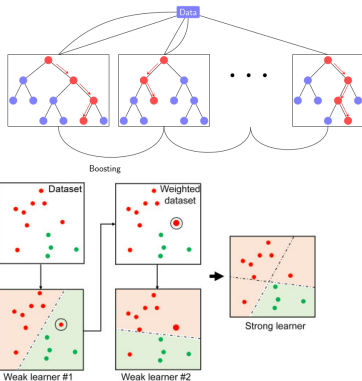voting
prediction

1. Create different (simple) tree models (stumps)
2. Each model is created with a subset of observation/features ($\sim$2m/3)
3. We combine the prediction of all trees

Main Parameters
- ▶ n_estimators: Number of trees
- ▶ max_features: Number of features selected for the split
- ▶ bootstrap=False: Use all samples
- ▶ Tree parameters

ADABOOST

1. Assign equal weights to observations $w_i^{(0)} = 1/m$

2. For $k = 1, \ldots, K$
   - Select a sample of observations based on the weights.
   - Create the k-th weak learner and compute predictions $x^{(k)}$
   - Compute the model **weighted** error and assign its coefficient:

   $$\alpha^{(k)} = \lambda \times \log((1 - error)/error)$$

   - Update sample weights:

   $$w_i^{(k+1)} \propto w_i^{(k)} \times \exp(-\alpha^{(k)} y_i \hat{x}_i^{(k)})$$

3. Final weighted prediction

Main Parameters

▶ n_estimators: Number of estimators (K)

▶ base_estimator: Weak estimator type

▶ learning_rate: weights of estimator in final decision (λ)

THANK YOU

POLIMI GRADUATE SCHOOL OF MANAGEMENT